

## Task 1 - Preparation

### 1.1. Choose an Application

For this task, I have chosen an open-source project called “Gitea”, a self-hosted Git service. It is lightweight, easy to deploy, and meets the requirements outlined in the task.

- Project Repository: [Gitea GitHub Repository](https://github.com/go-gitea/gitea)
- Brief Description: Gitea is a community-managed lightweight code hosting solution written in Go. It is published under the MIT license and can be run as a standalone binary or in a Docker container. It provides features similar to GitHub, such as repository management, issue tracking, pull requests, and more.
- Why Gitea?:
  - It uses configuration files in JSON, YAML, or TOML formats.
  - It supports environment variables for configuration.
  - It requires secrets for database credentials, OAuth tokens, etc.
  - It can be exposed to the internet (optional).
  - It has health check endpoints (e.g., `/api/v1/healthz`).

### 1.2. Get Familiar with Kubernetes (k8s) and Concepts

Kubernetes (k8s) is an open-source platform for automating deployment, scaling, and operations of application containers across clusters of hosts. Key concepts to understand include:

- **Pods**: The smallest deployable units in Kubernetes, which can contain one or more containers.
- **Services**: An abstraction that defines a logical set of Pods and a policy to access them.
- **Deployments**: Manage the deployment and scaling of Pods.
- **ConfigMaps and Secrets**: Used to manage configuration data and sensitive information, respectively.
- **Ingress**: Manages external access to services in a cluster.
- **Namespaces**: Provide a scope for resources within a cluster.

### 1.3. Install and Set Up the Necessary Tools

To work with Kubernetes, you need to install the following tools:

#### 1. **kubectl** :

- The Kubernetes command-line tool used to interact with the Kubernetes cluster.
- Installation instructions:
  - For Linux:

```
curl -LO "https://dl.k8s.io/release/${curl -L -s https://dl.k8s.io/release/stable.txt}/bin/linux/amd64/kubectl"
sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
```

```

• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps$ kubectl
kubectl controls the Kubernetes cluster manager.

Find more information at: https://kubernetes.io/docs/reference/kubectl/

Basic Commands (Beginner):
  create      Create a resource from a file or from stdin
  expose      Take a replication controller, service, deployment or pod and
  expose it as a new Kubernetes service
  run         Run a particular image on the cluster
  set         Set specific features on objects

```

## 2. Minikube :

- A tool that allows you to run a single-node Kubernetes cluster locally.
- Installation instructions:
  - For Linux:

```

curl -LO
https://storage.googleapis.com/minikube/releases/latest/minikube-linux-a
md64
sudo install minikube-linux-amd64 /usr/local/bin/minikube

```

### 3. Start Minikube:

- After installation, start Minikube:

```
minikube start
```

```

• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps$ minikube start
🐸 minikube v1.35.0 on Ubuntu 24.04
🌟 Using the docker driver based on existing profile
👍 Starting "minikube" primary control-plane node in "minikube" cluster
🚚 Pulling base image v0.0.46 ...
🔄 Updating the running docker "minikube" container ...

```

- Verify the cluster is running:

```
kubectl get nodes
```

### 1.4. Get Access to Kubernetes Dashboard

The Kubernetes Dashboard provides a web-based UI for managing your cluster.

#### 1. Enable the Dashboard :

- Minikube includes the Kubernetes Dashboard as an addon. Enable it with:

```
minikube addons enable dashboard
```

#### 2. Access the Dashboard :

- Start the dashboard proxy:

## kubectl proxy

- Access the dashboard in your browser:

```
http://127.0.0.1:35247/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:/proxy/#/workloads?namespace=default
```

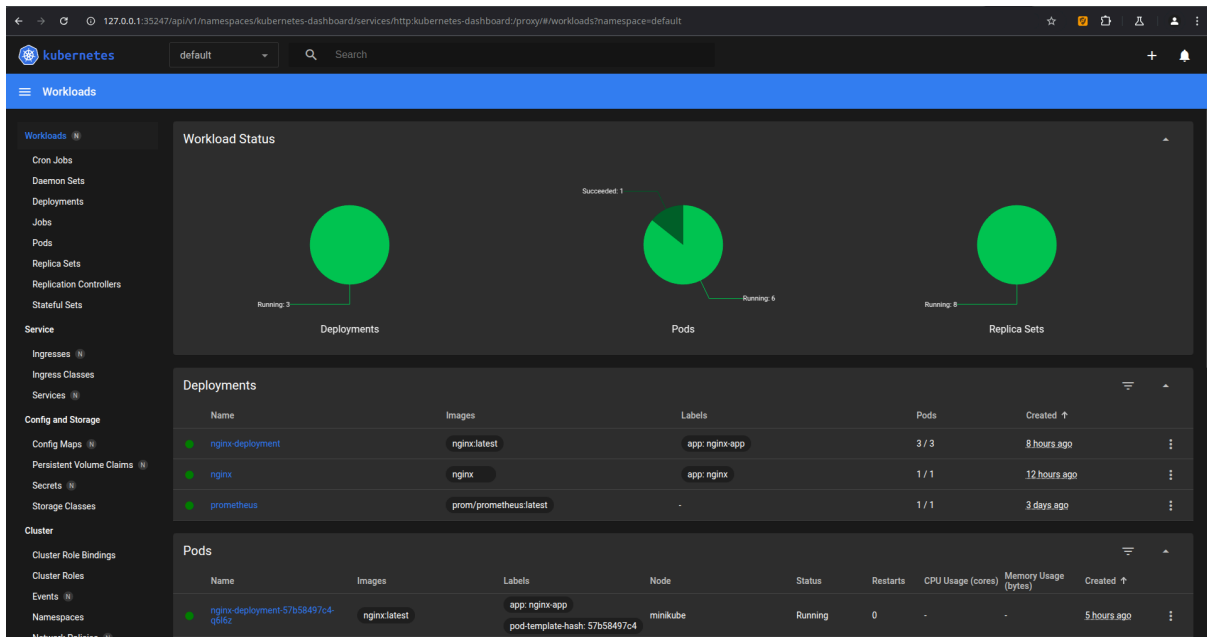
### 3. Authenticate :

- You can use the token from the default service account to log in:

```
kubectl -n kubernetes-dashboard create token admin-user
```

- Alternatively, use the Minikube token:

```
minikube dashboard
```



---

## Task 2 - k8s Nodes

### Étapes :

1.

Start a cluster local with Minikube :

```
minikube start
```

○

2. Listen and describe the nodes of cluster

lists of nodes of my cluster :

```
kubectl get nodes
```

```
● etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl
get nodes
NAME          STATUS    ROLES          AGE      VERSION
minikube      Ready    control-plane  5d17h    v1.32.0
○ etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ █
```

get more detail of one node

```
kubectl describe node minikube
```

```
● etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl describe node minikube
Name:          minikube
Roles:         control-plane
Labels:        beta.kubernetes.io/arch=amd64
               beta.kubernetes.io/os=linux
               kubernetes.io/arch=amd64
               kubernetes.io/hostname=minikube
               kubernetes.io/os=linux
               minikube.k8s.io/commit=dd5d320e41b5451cdf3c01891bc4e13d189586ed-dirty
               minikube.k8s.io/name=minikube
               minikube.k8s.io/primary=true
               minikube.k8s.io/updated_at=2025_02_18T15_31_03_0700
               minikube.k8s.io/version=v1.35.0
               node-role.kubernetes.io/control-plane=
               node.kubernetes.io/exclude-from-external-load-balancers=
Annotations:   kubeadm.alpha.kubernetes.io/cri-socket: unix:///var/run/cni-dockerd.sock
               node.alpha.kubernetes.io/ttl: 0
               volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Tue, 18 Feb 2025 15:30:53 +0300
Taints:         <none>
Unschedulable:  false
Lease:
  HolderIdentity: minikube
  AcquireTime:    <unset>
```

### 3. Get more system information about a node

get info about OS , CPU of node :

```
kubectl get nodes -o wide
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl get nodes -o wide
NAME          STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION
minikube      Ready     control-plane  5d18h  v1.32.0   192.168.49.2   <none>        Ubuntu 22.04.5 LTS   6.8.0-52-generic
docker://27.4.1
```

check détails of OS and ressources CPU :

```
kubectl describe node minikube | grep -i "os\|kernel\|cpu"
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl describe node minikube | g
rep -i "os\|kernel\|cpu"
      beta.kubernetes.io/os=linux
      kubernetes.io/hostname=minikube
      kubernetes.io/os=linux
Ready               True    Mon, 24 Feb 2025 09:29:54 +0300   Tue, 18 Feb 2025 15:30:56 +0300   KubeletReady
  kubelet is posting ready status
  Hostname:      minikube
  cpu:           4
  cpu:           4
  Kernel Version: 6.8.0-52-generic
  OS Image:      Ubuntu 22.04.5 LTS
  Namespace      Name
  mory Limits   Age
  cpu           3050m (76%)  2300m (57%)
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$
```

---

## Task 3 - k8s Pod

### Objectif :

Create and manipulate a Kubernetes Pod that runs an application.

### Steps :

#### 1. Understand Structure of Pod

- A Pod is a basic unit in Kubernetes. It contains one or more containers.
- The specifications of a Pod are defined in a YAML file with the main fields :
  - **apiVersion**: Kubernetes API version
  - **kind**: resource type (here, Pod)
  - **metadata**: information about the Pod (name, labels, etc.)
  - **spec**: specifications, including containers and their parameters

Here is an example of a Pod running an Nginx container:

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx-container
    image: nginx
    ports:
    - containerPort: 80
```

Create this file, for example pod.yaml, and then apply it:

```
kubectl apply -f pod.yaml
```

### Check Pod Deployment

List active Pods:

```
kubectl get pods
```

```

cpu 3050m (70%) 2500m (57%)
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
busybox       0/1     Completed 0           2d21h
nginx         1/1     Running   2 (3m4s ago) 2d20h

```

get more détails of Pod :

```
kubectl describe pod nginx
```

```

• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl describe pod nginx
Name:          nginx
Namespace:     default
Priority:       0
Service Account: default
Node:          minikube/192.168.49.2
Start Time:    Fri, 21 Feb 2025 12:54:06 +0300
Labels:        name=nginx-app
Annotations:    <none>
Status:        Running
IP:            10.244.0.43
IPs:
  IP: 10.244.0.43
Containers:
  nginx:
    Container ID:  docker://6ff2be162f1aa49ef9cab89ddb56eed03c864a5ed4acfe7d19b383a45500ac66
    Image:         nginx:latest
    Image ID:      docker-pullable://nginx@sha256:91734281c0ebfc6f1aea979cfeed5079cfe786228a71cc6f1f46a228cde6e34
    Port:         60/TCP
    Host Port:    0/TCP
    Command:
      /bin/sh
      -c
    Args:
      apt update && apt install -y iputils-ping && nginx -g 'daemon off;';
    State:        Running
      Started:    Mon, 24 Feb 2025 09:30:51 +0300
    Last State:   Terminated
      Reason:     Error
      Exit Code:  100
      Started:    Mon, 24 Feb 2025 09:28:54 +0300
      Finished:   Mon, 24 Feb 2025 09:30:25 +0300
    Ready:        True
    Restart Count: 2
    Limits:
      cpu:        500m
      memory:     128Mi
    Requests:
      cpu:        500m
      memory:     128Mi

```

Look at the container logs:

```
kubectl logs nginx
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl logs nginx
```

```
WARNING: apt does not have a stable CLI interface. Use with caution in scripts.
```

```
Get:1 http://deb.debian.org/debian bookworm InRelease [151 kB]
Get:2 http://deb.debian.org/debian bookworm-updates InRelease [55.4 kB]
Get:3 http://deb.debian.org/debian-security bookworm-security InRelease [48.0 kB]
Get:4 http://deb.debian.org/debian bookworm/main amd64 Packages [8792 kB]
Get:5 http://deb.debian.org/debian bookworm-updates/main amd64 Packages [13.5 kB]
Get:6 http://deb.debian.org/debian-security bookworm-security/main amd64 Packages [246 kB]
Fetched 9306 kB in 5s (1814 kB/s)
Reading package lists...
Building dependency tree...
Reading state information...
2 packages can be upgraded. Run 'apt list --upgradable' to see them.
```

Open shell inside a Pod :

```
kubectl exec -it nginx -- /bin/sh
```

Check if Nginx work :

```
curl http://localhost
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl exec -it nginx -- /bin/sh
# curl http://localhost
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
#
```



## Task 4 - k8s Service

A Service in Kubernetes allows to :

- Provide a stable IP to access Pods.
- Enable internal/external communication.
- Offer a load balancing mechanism.

### Key fields in a YAML specification file of a Service :

- **apiVersion**: Version of the Kubernetes API.
  - **kind**: Resource type (here Service).
  - **metadata**: Name of the Service.
  - **spec**:
    - **selector**: Labels used to identify the targeted Pods.
    - **ports**: Exposed ports (targetPort, port, nodePort if needed).
    - **type**: Service type (ClusterIP, NodePort, LoadBalancer).
- 

## 2. Write a Service for your Pod and deploy it

Assume that your Pod uses an Nginx container.

### Service Definition (**service.yaml**)

This Service exposes the application via an internal IP with ClusterIP :

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: my-nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: ClusterIP
```

## Déploiement of Service

Apply the file with :

```
kubectl apply -f service.yaml
```

Check that the Service is created :

```
kubectl get services
kubectl describe service nginx-service
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl get services
NAME                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes           ClusterIP   10.96.0.1    <none>        443/TCP          5d18h
my-nginx-external    NodePort    10.109.224.41 <none>        80:30080/TCP     2d20h
nginx-service        NodePort    10.98.120.249 <none>        80:30090/TCP     2d21h

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl describe service nginx-service
Name:                 nginx-service
Namespace:            default
Labels:               <none>
Annotations:          <none>
Selector:             app=nginx
Type:                 NodePort
IP Family Policy:     SingleStack
IP Families:          IPv4
IP:                   10.98.120.249
IPs:                  10.98.120.249
Port:                 <unset> 80/TCP
TargetPort:           80/TCP
NodePort:             <unset> 30090/TCP
Endpoints:            10.244.0.45:80
Session Affinity:     None
External Traffic Policy: Cluster
Internal Traffic Policy: Cluster
Events:               <none>
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$
```

## 3. Check communication between Pods

List Pods and their IP addresses :

```
kubectl get pods -o wide
```

```

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE
READINESS GATES
busybox                             0/1     Completed 0           2d21h  10.244.0.12   minikube   <none>
nginx                                1/1     Running   2 (8m8s ago)  2d20h  10.244.0.43   minikube   <none>
nginx-5869d7778c-td9jp              1/1     Running   1 (11m ago)  2d21h  10.244.0.45   minikube   <none>
nginx-deployment-57b58497c4-nw2hc    1/1     Running   1 (11m ago)  2d14h  10.244.0.37   minikube   <none>
nginx-deployment-57b58497c4-q6l6z    1/1     Running   1 (11m ago)  2d14h  10.244.0.47   minikube   <none>
nginx-deployment-57b58497c4-x79pn    1/1     Running   1 (11m ago)  2d14h  10.244.0.41   minikube   <none>
prometheus-76f97bf89f-5n5rc         1/1     Running   2 (11m ago)  5d16h  10.244.0.42   minikube   <none>

```

Each Pod has an IP assigned by Kubernetes. You can test the DNS resolution between them :

```
kubectl exec -it nginx -- ping nginx-service
```

```

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl exec -it nginx -- ping nginx-service
PING nginx-service.default.svc.cluster.local (10.98.120.249) 56(84) bytes of data.

```

If everything works fine, the Pod should be able to reach the Service by its DNS name.

## 4. Deleting and recreating a Pod

delete the Pod :

```
kubectl delete pod nginx
```

```

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl delete pod nginx
pod "nginx" deleted

```

Kubernetes will automatically recreate a new Pod if a Deployment is used. Check the IP of the new Pod and make sure it can still be reached by the Service :

```
kubectl get pods -o wide
```

```

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task4$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE       NOMINATED NODE
EADINESS GATES
busybox                             0/1     Completed 0           2d21h 10.244.0.12     minikube   <none>
none>
nginx                               1/1     Running   0           109s   10.244.0.48     minikube   <none>
none>
nginx-5869d7778c-td9jp              1/1     Running   1 (18m ago) 2d21h   10.244.0.45     minikube   <none>
none>
nginx-deployment-57b58497c4-nw2hc    1/1     Running   1 (17m ago) 2d14h   10.244.0.37     minikube   <none>
none>
nginx-deployment-57b58497c4-q6l6z    1/1     Running   1 (17m ago) 2d14h   10.244.0.47     minikube   <none>
none>
nginx-deployment-57b58497c4-x79pn    1/1     Running   1 (17m ago) 2d14h   10.244.0.41     minikube   <none>
none>
prometheus-76f97bf89f-5n5rc         1/1     Running   2 (18m ago) 5d16h   10.244.0.42     minikube   <none>
none>

```

The Pod IP will have changed, but the connection via the Service should still work.

## 5. Learn about LoadBalancer and NodePort

Kubernetes offers several types of Services :

- **ClusterIP** (défaut) : Accessible only from the cluster
- **NodePort** Exposes the Service on a specific port of each node.
- **LoadBalancer** : Creates an external Load Balancer (requires a cloud provider like AWS, GCP, Azure).

You can test a NodePort by modifying `service.yaml` :

type: NodePort

Apply the changes :

```

kubectl apply -f service.yaml
kubectl get services

```

```

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task4$ kubectl get services
NAME                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes          ClusterIP   10.96.0.1    <none>        443/TCP          5d18h
my-nginx-external    NodePort    10.109.224.41 <none>        80:30080/TCP     2d20h
nginx-service        NodePort    10.98.120.249 <none>        80:30090/TCP     2d21h

```

The assigned port will be in the range **30000-32767**.

You can access it via `http://10.98.120.249:30090`.

## 6. Deploy an External Service to access from the local host

If you want to test access from your computer:

```
apiVersion: v1
kind: Service
metadata:
  name: my-nginx-external
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      nodePort: 30080
  type: NodePort
```

```
kubectl apply -f service.yaml
minikube service my-nginx-external --url
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task4$ kubectl get services
NAME                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes           ClusterIP   10.96.0.1    <none>         443/TCP          5d18h
my-nginx-external    NodePort    10.109.224.41 <none>         80:30080/TCP     2d20h

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task4$ minikube service my-nginx-external
--url
http://192.168.49.2:30080
```

---

## Bonus: Labels et Sélecteurs

### Labels in Kubernetes

Labels are tags attached to Kubernetes resources (Pods, Services, Deployments, etc.).

Example of a Pod with a label:

```
metadata:
  labels:
    app: nginx
```

The **Service** selects the Pods that have this label:

```
selector:
  app: nginx
```

## Test 1: No Pod Matches the Label

If a **Service** has this selector:

```
selector:  
  app: missing-label
```

But no Pod has this label, then the command:

```
kubectl get endpoints my-nginx-service
```

Will return **0 endpoints**, meaning that no Pod is accessible through this Service.

## Test 2: Multiple Pods Match the Label

If you have multiple Pods with the same label:

```
kubectl scale deployment nginx --replicas=3  
kubectl get endpoints nginx-service
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task4$ kubectl scale deployment nginx --re  
plicas=3  
deployment.apps/nginx scaled  
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task4$ kubectl get endpoints nginx-service  
NAME                ENDPOINTS                                     AGE  
nginx-service        10.244.0.45:80,10.244.0.49:80,10.244.0.50:80  2d21h
```

The **Service** will automatically distribute traffic between these Pods. 🚀

# Task 5

---

## 1. Identifying the Required Fields in a Deployment Spec

The key fields in a **Deployment** manifest include :

- **apiVersion** : For example, `apps/v1`
- **kind** : Must be `Deployment`
- **metadata** : Contains the name, labels, etc.
- **spec** :
  - **replicas** : Number of desired Pods.
  - **selector** : Criteria (labels) to associate the Deployment with the Pods.
  - **template** : The Pod template that will be created, which itself consists of :
    - **metadata** (labels, annotations)
    - **spec** (list of containers, ports, volumes, environment variables, etc.)
  - *(Optional)* **strategy** : Defines the update strategy (e.g., RollingUpdate) with parameters like maxSurge and maxUnavailable..
  - *(Optional)* **revisionHistoryLimit** : Number of old versions to retain for rollback.

minimal YAML Deployment :

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
  labels:
    app: my-app
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-app-container
          image: nginx:latest
          ports:
            - containerPort: 80
```

---

## 2. Clean Up Previous Pod Manifests and Deploy the Application

Before creating the **Deployment**, make sure to delete or clean up any old Pods or manifests (created via `kubectl run` or individual YAML manifests).

Then, create a YAML file (e.g., `deployment.yaml`) with the Deployment manifest adapted to your application.

Apply it using:

```
kubectl apply -f deployment.yaml
```

## 3. List and Describe Deployments

To verify that the **Deployment** has been created and check its status, use:

```
kubectl get deployments
kubectl describe deployment my-app
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ touch deployment.yaml
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl apply -f deployment.yaml
deployment.apps/my-app created
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
my-app              2/2     2            2           12s
nginx               3/3     3            3           2d21h
nginx-deployment    3/3     3            3           2d17h
prometheus          1/1     1            1           5d17h
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl describe deployment my-app
Name:                my-app
Namespace:           default
CreationTimestamp:    Mon, 24 Feb 2025 10:01:40 +0300
Labels:               app=my-app
Annotations:          deployment.kubernetes.io/revision: 1
Selector:             app=my-app
Replicas:             2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType:         RollingUpdate
MinReadySeconds:      0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=my-app
```

The describe command provides details about the **ReplicaSet**, recorded events (useful for debugging), and update history.

## 4. Scale the Deployment to Three Replicas

To increase the number of Pods to **3**, you can either edit the YAML file:

```
spec:
  replicas: 3
```

Then apply the update:



```
kubectl apply -f deployment.yaml
```

Or use the command:

```
kubectl scale deployment my-app --replicas=3
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl scale deployment my-app --replicas=3
deployment.apps/my-app scaled
```

## 5. Access the Shell and Logs of Pods via Deployment Labels

The Pods created by the Deployment inherit the **labels** defined in the template.

- **List Pods with a specific label:**

```
kubectl get pods -l app=my-app
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl get pods -l app=my-app
NAME                                READY   STATUS    RESTARTS   AGE
my-app-6fd49669c9-4sb6z            1/1     Running   0           117s
my-app-6fd49669c9-bchzh            1/1     Running   0           117s
my-app-6fd49669c9-p2vxv            1/1     Running   0           35s
```

- **Access a Pod's shell (if the image supports it):**

```
kubectl exec -it my-app-6fd49669c9-4sb6z -- /bin/sh
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl exec -it my-app-6fd49669c9-4sb6z -- /bin/sh
# echo "test"
test
```

- **View logs of a specific Pod:**

```
kubectl logs my-app-6fd49669c9-4sb6z
```

```

• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl logs my-app-6fd49669c9-4sb6
z
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/02/24 07:01:45 [notice] 1#1: using the "epoll" event method
2025/02/24 07:01:45 [notice] 1#1: nginx/1.27.4
2025/02/24 07:01:45 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2025/02/24 07:01:45 [notice] 1#1: OS: Linux 6.8.0-52-generic
2025/02/24 07:01:45 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2025/02/24 07:01:45 [notice] 1#1: start worker processes
2025/02/24 07:01:45 [notice] 1#1: start worker process 29
2025/02/24 07:01:45 [notice] 1#1: start worker process 30
2025/02/24 07:01:45 [notice] 1#1: start worker process 31
2025/02/24 07:01:45 [notice] 1#1: start worker process 32
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ 

```

- View logs using labels (e.g., for all matching Pods):

```
kubectl logs -l app=my-app --tail=50
```

```

• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl logs -l app=my-app --tail=5
0
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/02/24 07:01:45 [notice] 1#1: using the "epoll" event method
2025/02/24 07:01:45 [notice] 1#1: nginx/1.27.4
2025/02/24 07:01:45 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2025/02/24 07:01:45 [notice] 1#1: OS: Linux 6.8.0-52-generic
2025/02/24 07:01:45 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2025/02/24 07:01:45 [notice] 1#1: start worker processes
2025/02/24 07:01:45 [notice] 1#1: start worker process 29
2025/02/24 07:01:45 [notice] 1#1: start worker process 30
2025/02/24 07:01:45 [notice] 1#1: start worker process 31

```

## 6. Update the Application Configuration and Observe the Behavior

Modify part of the **Deployment manifest** (e.g., update the image version or environment variables).

For example, change the image:

image: my-app-image:2.0

Then apply the changes:

```
kubectl apply -f deployment.yaml
```

Use `--watch` to follow the update process:

```
kubectl get pods --watch
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl get pods --watch
NAME                                READY   STATUS    RESTARTS   AGE
busybox                             0/1     Completed 0           2d22h
my-app-6fd49669c9-4sb6z             1/1     Running   0           13m
my-app-6fd49669c9-bchzh             1/1     Running   0           13m
my-app-6fd49669c9-p2vxv             1/1     Running   0           12m
nginx                                1/1     Running   0           32m
nginx-5869d7778c-k6gpz              1/1     Running   0           20m
nginx-5869d7778c-q6vsr              1/1     Running   0           20m
nginx-5869d7778c-td9jp              1/1     Running   1 (48m ago) 2d21h
nginx-deployment-57b58497c4-nw2hc    1/1     Running   1 (48m ago) 2d14h
nginx-deployment-57b58497c4-q6l6z    1/1     Running   1 (48m ago) 2d14h
nginx-deployment-57b58497c4-x79pn    1/1     Running   1 (48m ago) 2d14h
prometheus-76f97bf89f-5n5rc         1/1     Running   2 (48m ago) 5d17h
```

You will see Kubernetes perform a **rolling update**, gradually replacing old Pods with new ones.

## 7. Perform a Rollback to the Previous Version

If the update causes issues, rollback to the previous version with:

```
kubectl rollout undo deployment my-app
```

To check the update history:

```
kubectl rollout history deployment my-app
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl rollout history deployment my-app
deployment.apps/my-app
REVISION  CHANGE-CAUSE
1          <none>
```

## 8. Define CPU and Memory Requests & Limits

Set resource constraints in the **containers** section of the Deployment template:

```
resources:
  requests:
    cpu: "100m"
    memory: "128Mi"
  limits:
    cpu: "500m"
    memory: "256Mi"
```

### Behavior When Limits Are Exceeded

- **CPU Throttling:** If the app uses more than 500m (0.5 CPU), it will be slowed down.

- **Memory OOM Kill:** If the app exceeds 256Mi, Kubernetes will terminate the Pod (OOMKilled) and restart it based on the restart policy.

To test this, generate a high load on the app and observe **throttling** or **OOM kills** using:

`kubectl top pods`

## Task 6

---

### 1. Fields needed in a Secret specification

A Secret manifest contains the following essential fields:

- **apiVersion** : Usually `v1`.
  - **kind** : Must be `Secret`.
  - **metadata** : Name, labels, and annotations.
  - **type** : The type of the secret (often `Opaque` for a generic secret).
  - **data** : A dictionary of key/value pairs where each value must be Base64 encoded.  
→ Alternatively, you can use `stringData` which allows you to specify the values in plain text; Kubernetes will take care of the encoding.
- 

### 2. Creating and Applying a Secret Manifest

For example, let's create a `secret.yaml` file to store a username and password:

```
apiVersion: v1
kind: Secret
metadata:
  name: my-secret
type: Opaque
data:
  # "myuser" encodé en base64 : echo -n "myuser" | base64 → bX11c2Vy
  username: bX11c2Vy
  # "mypassword" encodé en base64 : echo -n "mypassword" | base64 →
  bXlwYXNzd29yZA==
  password: bXlwYXNzd29yZA==
```

You can also use the `stringData` section to avoid doing the encoding manually:

```
apiVersion: v1
kind: Secret
metadata:
  name: my-secret
type: Opaque
stringData:
  username: myuser
  password: mypassword
```

apply secret to the cluster :

```
kubectl apply -f secret.yaml
```

---

### 3. Retrieve and describe the Secret with kubectl

List your secrets:

```
kubectl get secrets
```

```
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl get secrets
NAME          TYPE      DATA   AGE
json-secret   Opaque    1       2d16h
my-secret     Opaque    2       2d16h
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$
```

describe secret

```
kubectl describe secret my-secret
```

```
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ kubectl describe secret my-secret
Name:         my-secret
Namespace:    default
Labels:       <none>
Annotations:  <none>

Type: Opaque

Data
====
password: 10 bytes
username: 6 bytes
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$
```

---

### 4. Decode Secrets

To check the value of a field, you can copy the Base64 value and decode it locally. For example:

```
echo 'bXl1c2Vy' | base64 --decode
```

```
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$ echo 'bXl1c2Vy' | base64 --decode
myuser
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task5$
```

---

## 5. Update your Deployment to reference the Secret as an environment variable

Modify your Deployment manifest to inject the Secret data into the containers. For example, in your deployment.yaml file:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx-app
spec:
  replicas: 3 # Deploy 2 instances of the pod
  selector:
    matchLabels:
      app: nginx-app # Ensures Deployment manages the right Pods
  template:
    metadata:
      labels:
        app: nginx-app
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          command: ["/bin/sh", "-c"]
          args:
            - apt update && apt install -y iputils-ping net-tools &&
nginx -g 'daemon off;';
          resources:
            limits:
              memory: "128Mi"
              cpu: "600m"
          ports:
            - containerPort: 60
          env:
            - name: APP_USERNAME
              valueFrom:
                secretKeyRef:
                  name: my-secret
```

```
        key: username
-   name: APP_PASSWORD
    valueFrom:
      secretKeyRef:
        name: my-secret
        key: password
```

Apply the update :

```
kubectl apply -f secret-deployment.yaml
```

---

## 6. Verify that the Secret is available in the Pod

You can run a command in one of the pods to display the environment variables:

List the pods with the Deployment label:

```
kubectl get pods -l app=nginx-app
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl get pods -l app=nginx-app
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-d45596dc8-828xc    1/1     Running   0           80s
nginx-deployment-d45596dc8-gf69f    1/1     Running   0          117s
nginx-deployment-d45596dc8-rs5lk    1/1     Running   0           2m2s
```

go to the shell of pod

```
kubectl exec -it nginx-deployment-d45596dc8-rs5lk -- /bin/sh
```

check the variable :

```
echo $APP_USERNAME
echo $APP_PASSWORD
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl exec -it nginx-deployment-d45596dc8-rs5lk -- /bin/sh
# echo $APP_USERNAME
echo $APP_PASSWORDmyuser
#
mypassword
#
```



---

## Bonus: Creating a Secret from a JSON File

Suppose you have a file named `credentials.json` containing:

```
{
  "username": "jsonuser",
  "password": "jsonpass"
}
```

You can create a **Secret** from this file using a similar approach to ConfigMaps (but specifically for secrets):

```
kubectl create secret generic json-secret --from-file=credentials.json
```

This command creates a **secret** named `json-secret` with a key `credentials.json`, where the value is the file's content encoded in **Base64**.

## Verifying the Secret

To check the created secret:

```
kubectl get secret json-secret -o yaml
kubectl describe secret json-secret
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl get secret json-secret -o y
aml
apiVersion: v1
data:
  credentials.json: ewogICAgInVzZXJuYV1lIjogImpzb251c2VyIiwKICAgICJwYXNkd29yZCI6ICJqc29ucGFzcycKICB9CiAg
kind: Secret
metadata:
  creationTimestamp: "2025-02-21T15:06:20Z"
  name: json-secret
  namespace: default
  resourceVersion: "102754"
  uid: 13c20125-2af0-4f8e-8009-3f8c08e9faa9
type: Opaque
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl describe secret json-secret
Name:         json-secret
Namespace:    default
Labels:       <none>
Annotations:  <none>

Type:  Opaque

Data
====
credentials.json:  63 bytes
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$
```

To **decode** the stored data, retrieve the Base64 value and use:

```
echo 'encoded_value' | base64 --decode
```

## Injecting Secret Values into a Deployment

There are two ways to inject these secrets into your **Deployment**:

### 1 Direct Injection into Environment Variables

If you want to extract specific keys from the JSON file, it's better to create the **secret** with multiple key-value pairs. You can use a tool like `jq` to extract values:

```
kubectl create secret generic json-secret \
  --from-literal=username=$(jq -r '.username' credentials.json) \
  --from-literal=password=$(jq -r '.password' credentials.json)
```

Then, reference them in your **Deployment** as environment variables.

### 2 Mounting the Secret as a Volume

You can mount the **Secret** as a volume inside the Pod so that the JSON file is accessible from within the container.

 **This allows Kubernetes applications to securely manage sensitive data like credentials!**

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
spec:
  replicas: 1
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-app-container
          image: my-app-image:1.0
          volumeMounts:
            - name: json-secret-volume
```

```
    mountPath: /etc/secrets
    readOnly: true
volumes:
- name: json-secret-volume
  secret:
    secretName: json-secret
```

1. In this case, the /etc/secrets/credentials.json file will be available in the pod and you can use it in your application.

## Task 7

---

### 1. The Required Fields in a ConfigMap Manifest

The essential elements of a **ConfigMap** manifest are:

- **apiVersion**: usually v1
- **kind**: must be ConfigMap
- **metadata**: includes the name, labels, and optional annotations
- **data**: a dictionary of key-value pairs

→ You can also use **stringData** to provide unencoded values (Kubernetes will handle the Base64 conversion).

#### Minimal Example:

*(Example YAML not provided in the text.)*

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-app-config
data:
  APP_ENV: "production"
  APP_DEBUG: "false"
```

---

### 2. Modifying the Deployment to Inject Environment Variables

To allow your application to receive this configuration via **environment variables**, update the **Deployment** manifest by adding a reference to the **ConfigMap**.

For example, in the **envFrom** section of the container:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
```

```

labels:
  app: my-app
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-app-container
          image: my-app-image:1.0
          ports:
            - containerPort: 80
          envFrom:
            - configMapRef:
                name: my-app-config

```

### 3. Creating a ConfigMap with Key/Value Pairs

Create a YAML file (e.g., configmap.yaml) that contains the desired configuration data:

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: my-app-config
data:
  APP_ENV: "production"
  APP_DEBUG: "false"

```

Apply it using:

```
kubectl apply -f configmap.yaml
```

```

● etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl get configmap
NAME                DATA  AGE
kube-root-ca.crt    1      5d19h
my-app-config        2      2d15h
my-app-config-json   1      2d15h

```

Update your **Deployment** (as shown in step 2) to use this **ConfigMap**.

#### 4. Creating a JSON Configuration File

Create a file named, for example, config.json, containing your JSON data:

```
{
  "log_level": "info",
  "max_connections": 100
}
```

---

#### 5. Creating a ConfigMap from a JSON File

You have two options:

##### Option A – Directly in the YAML manifest:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-app-config-json
data:
  config.json: |
    {
      "log_level": "info",
      "max_connections": 100
    }
```

##### Option B – Using the Command Line:

```
kubectl create configmap my-app-config-json --from-file=config.json
```

#### 6. Updating the Deployment to Mount the ConfigMap as a Volume

To make the JSON file accessible within the container's filesystem, modify your **Deployment** to add a **volume** and a **volumeMount**:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
```

```
labels:
  app: my-app
spec:
  replicas: 2
```

```
selector:
  matchLabels:
    app: nginx-app
template:
  metadata:
    labels:
      app: nginx-app
  spec:
    containers:
      - name: my-app-container
        image: my-app-image:1.0
        ports:
          - containerPort: 80
        envFrom:
          - configMapRef:
              name: my-app-config
        volumeMounts:
          - name: config-volume
            mountPath: /etc/config # Le chemin où sera monté le fichier
JSON
    volumes:
      - name: config-volume
        configMap:
          name: my-app-config-json
```

apply the changes with :

```
kubectl apply -f deployment.yaml
```

---

## 7. Check ConfigMap with kubectl

To view the contents of the ConfigMap and verify that it is stored in clear text, use

```
kubectl get configmap my-app-config -o yaml
kubectl describe configmap my-app-config
```

```

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl get configmap my-app-config -o yaml
apiVersion: v1
data:
  APP_DEBUG: "false"
  APP_ENV: production
kind: ConfigMap
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","data":{"APP_DEBUG":"false","APP_ENV":"production"},"kind":"ConfigMap","metadata":{"annotations":{},"name":"my-app-config","namespace":"default"}}
  creationTimestamp: "2025-02-21T16:02:44Z"
  name: my-app-config
  namespace: default
  resourceVersion: "105281"
  uid: 72474e7b-9139-4677-8ab5-111aaf925c73

```

```

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl describe configmap my-app-config
Name:         my-app-config
Namespace:    default
Labels:       <none>
Annotations:  <none>

Data
====
APP_DEBUG:
----
false

APP_ENV:
----
production

BinaryData
====

Events: <none>
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$

```

## 8. Verification in the container

To verify that the JSON file has been mounted, select a pod and open a shell:

```

kubectl get pods -l app=nginx-app
kubectl exec -it <pod-name> -- /bin/sh

```

Once in the shell, display the contents of the mounted file:

```

cat /etc/config/config.json

```



Events: <none>

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task7\$ kubectl get pods -l app=nginx-app

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-57b58497c4-nw2hc	1/1	Running	0	78s
nginx-deployment-57b58497c4-q6l6z	1/1	Running	0	36s
nginx-deployment-57b58497c4-x79pn	1/1	Running	0	73s
nginx-deployment-d45596dc8-wvg4r	1/1	Terminating	0	89m

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task7\$ kubectl exec -it nginx-deployment-57b58497c4-nw2hc -- /bin/sh

# cat /etc/config/config.json

```
{
  "log_level": "info",
  "max_connections": 100
}
#
```

Ln 14, Col 1: Spaces: 2, HT5: 8, LF: 1, YAMML: 0 @ Go Live: 00 kubernetecr://sche

## Task 8

---

### 1. Required Fields for a Namespace Manifest

A **Namespace** manifest is very simple. The essential elements are:

- **apiVersion:** v1
- **kind:** Namespace
- **metadata:** Contains at least the **name** field (optionally, labels or annotations).

#### Minimal Example:

*(Example YAML not provided in the text.)*

```
apiVersion: v1
kind: Namespace
metadata:
  name: app1-namespace
```

---

### 2. Creating Two Namespaces in the Cluster

You can create two separate YAML files or a single combined file. For example:

```
apiVersion: v1
kind: Namespace
metadata:
  name: app1-namespace
---
apiVersion: v1
kind: Namespace
metadata:
  name: app2-namespace
```

Apply the config :

```
kubectl apply -f namespaces.yaml
```

---

### 3. Retrieve and describe your Namespaces with kubectl

Lists of Specifics :

```
kubectl get namespaces
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl get namespaces
NAME                STATUS   AGE
app1-namespace      Active   2d15h
app2-namespace      Active   2d15h
```

Description of a specifics Namespace :

```
kubectl describe namespace app1-namespace
kubectl describe namespace app2-namespace
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl describe namespace app1-namespace
Name:          app1-namespace
Labels:        kubernetes.io/metadata.name=app1-namespace
Annotations:   <none>
Status:        Active

No resource quota.

No LimitRange resource.
Name:          app2-namespace
Labels:        kubernetes.io/metadata.name=app2-namespace
Annotations:   <none>
Status:        Active

No resource quota.

No LimitRange resource.
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$
```

#### 4. Deploying Two Different Applications in Separate Namespaces

You can deploy, for example, an **Nginx** application in app1-namespace and an **Apache** application in app2-namespace.

Here is an example of a combined manifest:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  namespace: app1-namespace
  labels:
    app: nginx
spec:
  replicas: 1
  selector:
```

```

    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apache-deployment
  namespace: app2-namespace
  labels:
    app: apache
spec:
  replicas: 1
  selector:
    matchLabels:
      app: apache
  template:
    metadata:
      labels:
        app: apache
    spec:
      containers:
        - name: apache
          image: httpd:latest
          ports:
            - containerPort: 80

```

apply the deployment :

```
kubectl apply -f deployments.yaml
```

---

## 5. Récupérer et décrire les Pods dans différents Namespaces

To list pods in the app1-namespace namespace:

```
kubectl get pods -n app1-namespace
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubectl get pods -n app1-namespace
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-96b9d695-rhgqh     1/1     Running   1 (81m ago)  2d15h

etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ kubectl describe pod -n app1-namespace nginx-deployment-96b9d695-rhgqh
Name:                               nginx-deployment-96b9d695-rhgqh
Namespace:                           app1-namespace
Priority:                             0
Service Account:                       default
Node:                                 minikube/192.168.49.2
Start Time:                           Fri, 21 Feb 2025 19:42:22 +0300
Labels:                               app=nginx
                                      pod-template-hash=96b9d695
Annotations:                           <none>
Status:                               Running
IP:                                   10.244.0.35
IPs:                                   <none>
```

## 6. Visibility and Connectivity Between Resources in Different Namespaces

By default, **Namespaces** isolate resources.

- You **cannot list** resources from another Namespace unless you use the `-n` flag or `--all-namespaces`.
- To connect two services in different **Namespaces**, you must use **fully qualified DNS names** (e.g., `service-name.namespace.svc.cluster.local`) or configure specific routing.

This means that, by default, an application in app1-namespace **will not automatically have access** to a resource (Pod, Service) in app2-namespace without explicit configuration.

---

### Bonus: Using kubectl/kubens

The tools `kubectx` and `kubens` allow you to quickly switch between contexts or namespaces without needing to specify the `-n` flag in each command.

You can install them from GitHub: [github.com/ahmetb/kubectx](https://github.com/ahmetb/kubectx).

```
See 'snap info kubectx' for additional versions.
• etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ sudo apt
  install kubectx
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  kubectx
0 upgraded, 1 newly installed, 0 to remove and 94 not upgraded.
Need to get 7,832 kB of archives.
After this operation, 34.2 MB of additional disk space will be used.
Get:1 http://ru.archive.ubuntu.com/ubuntu noble-updates/universe amd64 kubectx amd64 0.9.5-1ubu
ntu0.3 [7,832 kB]
Fetched 7,832 kB in 3s (2,800 kB/s)
Selecting previously unselected package kubectx.
(Reading database ... 499598 files and directories currently installed.)
Preparing to unpack .../kubectx_0.9.5-1ubuntu0.3_amd64.deb ...
Unpacking kubectx (0.9.5-1ubuntu0.3) ...
```

For example, to switch to the app1-namespace namespace:

```
kubens app1-namespace
```

```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$ kubens app1-namespace
✓Active namespace is "app1-namespace"
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task6$
```

Once the namespace is changed, you can simply run:

```
kubectl get pods
```

to view the pods in the current namespace.

Similarly, you can use kubectx to switch between contexts if you're managing multiple clusters.

### In Summary, This Task Allows You To:

- Create and isolate environments via namespaces.
- Deploy applications in different namespaces.
- View and monitor resources using the -n flag.
- Use tools like kubens to make multi-namespace management easier.

## Rancher for Centralized Kubernetes Cluster Management in Production

### Context and Objectives

In a complex production environment, having a centralized graphical interface allows you to:

- Quickly visualize the global state of the cluster (nodes, pods, services, deployments, etc.).
- Manage and control resources (adding/removing clusters, updating workloads, managing users).
- Monitor performance and metrics in real-time.
- Simplify multi-cluster management and security administration.

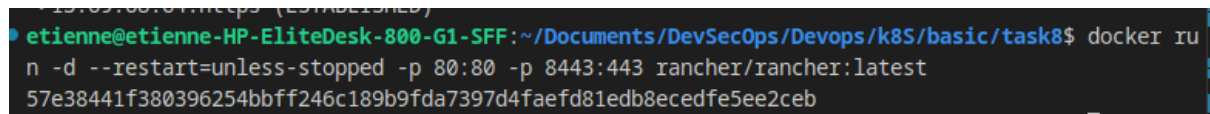
Rancher is one of the most popular and comprehensive solutions to meet these needs.

## 1. Rancher Installation

### Installation with Docker (Quick method for a test environment)

You can install Rancher by simply running a Docker container. For example, on a Linux machine:

```
docker run -d --restart=unless-stopped -p 80:80 -p 443:443
rancher/rancher:latest
```



```
etienne@etienne-HP-EliteDesk-800-G1-SFF:~/Documents/DevSecOps/Devops/k8S/basic/task8$ docker run -d --restart=unless-stopped -p 80:80 -p 8443:443 rancher/rancher:latest
57e38441f380396254bbff246c189b9fda7397d4faefd81edb8ecedfe5ee2ceb
```

### Production Installation

For a production setup, it's recommended to install Rancher on a Kubernetes cluster itself, for example, using Helm:

1. Add the Rancher Helm repository:

```
helm repo add rancher-latest
https://releases.rancher.com/server-charts/latest
helm repo update
```

2. Install Rancher in a dedicated namespace (e.g., cattle-system):

```
kubectl create namespace cattle-system
helm install rancher rancher-latest/rancher \
  --namespace cattle-system \
  --set hostname=rancher.mycompany.com
```

You'll need to configure DNS so the hostname points to the cluster's IP or use an Ingress Controller.

## 2. Access and Practical Actions via Rancher Interface

Once Rancher is installed, you can access the interface via your browser at an address like <https://rancher.mycompany.com>.

### Typical Actions Available in Rancher:

- **Add and Manage Clusters**

Rancher allows you to add new clusters (local or cloud) and group them into projects to better organize your resources.

- **Visualize and Control Resources**

You can view the real-time status of nodes, pods, services, and deployments. The interface provides dashboards with metrics (CPU, memory, etc.) and detailed logs.

- **Deploy Workloads**

From the Rancher interface, you can create deployments, StatefulSets, or jobs without typing kubectl commands. Simply fill out forms to define replicas, images, environment variables, etc.

- **Updates and Rollbacks**

Rancher provides the ability to easily update deployments and rollback to a previous version if needed. You can view the update history and initiate a rollback with a few clicks.

- **User and Role Management**

You can define access policies for users or teams, restricting access to specific clusters or projects.

- **Access to Shell Interface**

Rancher includes a web terminal that lets you run kubectl commands directly from the interface, making checks and troubleshooting easier.

- **Context Switching Tools**

Rancher simplifies multi-cluster management by clearly displaying which cluster or project is currently selected and allows you to switch between them easily.